

CompliVibe

THE AI TRUST INFRASTRUCTURE COMPANY

COMPANY & PRODUCT REPORT

The complete picture — positioning, product, traction, market, model, plans, and team, for leadership, partners, and investor conversations.

CompliVibe Private Limited · Incorporated June 2026 · Gurugram, India · Global operations
Confidential · Version 1.0

01 EXECUTIVE SUMMARY

CompliVibe is the trust layer for the AI era — the system organisations use to prove their AI is safe, compliant, and accountable, anywhere in the world.

AI is becoming foundational to every company. The bottleneck to deploying it is no longer capability — it is trust. Regulators, enterprise buyers, boards, and insurers all need proof that an AI system is governed. Today there is no default way to prove it. CompliVibe is building that default: a unified platform spanning AI inventory, risk, security, observability, evidence, and trust reporting across every major framework.

Unlike security-compliance incumbents that bolted AI on as an afterthought, CompliVibe is AI-native from day one — a global, enterprise-grade platform, not an MVP. It is live, with paying customers acquired within days of launch, patented technology, and a data room open to investors.

6

frameworks covered

828

mapped obligations

5

patents filed

48h

update velocity

02 POSITIONING

Not a compliance tool. The global trust infrastructure for AI.

Compliance is one outcome of trustworthy AI, not the destination. CompliVibe is positioned as infrastructure — analogous to what cloud became for software. Every organisation will eventually need both AI infrastructure and AI trust infrastructure; CompliVibe intends to be the latter, globally, from day one.

- **Unified** — A unified system replacing fragmented spreadsheets and manual governance.
- **AI-native** — AI-native, covering governance, risk, observability, evidence, and compliance automation in one platform.
- **Global** — Global and enterprise-grade — built for worldwide customers, not one region.

03 THE PROBLEM & WHY NOW

A regulatory forcing function meets a universal, unaddressed gap.

Across verticals, companies hold security certifications (SOC 2) but publish almost no AI-specific governance artifacts — model cards, bias testing, AI risk assessments. That is the universal gap. Meanwhile the EU AI Act's high-risk phase-in (August 2, 2026) is a hard, time-boxed deadline that forces action, and MiCA, the GENIUS Act, and DPDP raise the bar in parallel.

- **Timing** — Enforcement is accelerating worldwide, not slowing.
- **New risk** — AI adds new failure surfaces — insecure generated code, opaque data processing.
- **Demand** — Buyers, insurers, and regulators increasingly demand proof of trust up front.

04 THE PRODUCT

One platform to know, prove, and maintain AI trust.

CompliVibe unifies the full AI-trust lifecycle: discover the AI in use, assess it against applicable frameworks, generate the evidence, and keep it current as regulations change. The platform spans AI inventory, risk management, security, observability, evidence vaulting, and trust reporting.

Coverage & capability

- **Frameworks (6)** — EU AI Act, DPDP, ISO 42001, NIST AI RMF, SOC 2, and Colorado AI Act.
- **Depth** — 828 mapped obligations, including 651 EU AI Act obligations, with a 48-hour update velocity when rules change.
- **Shadow AI Discovery Engine** — a patented capability that surfaces unmanaged AI in an organisation.
- **Integrations** — 23+ integrations to pull evidence and signals from a customer's stack.

Feature-level specifics (Trust Score engine internals, evidence vault outputs, update-pipeline triggers) can be expanded in a dedicated product appendix.

05 THE TRUST SCORE & BENCHMARK

An objective, reproducible measure of how audit-ready an organisation's AI is.

The Trust Score maps a company's AI systems against all applicable obligations, marks each satisfied / partial / missing, and computes a risk-adjusted, weighted ratio into a 0–100 score, broken down by framework with a gap list ranked by deal-risk.

- **Benchmarked against** — regulatory ground truth (defensible today).
- **And against** — audit-readiness for a specific deadline (the sales trigger).
- **And eventually** — peer comparison as the anonymised dataset compounds (the moat).

06 TRACTION

Live, paying, and protected — built product-first, then raised.

Signal	Status
Product	MVP live, enterprise-grade
Paying customers	3 (acquired within days of launch)
Pipeline	13-account waitlist, 5 beta trials
Coverage	828 mapped obligations across 6 frameworks
Integrations	23+
IP	5 patents filed (incl. Shadow AI Discovery Engine)
Investor readiness	Live data room

This venture deliberately built product and won paying customers before seeking investment.

07 MARKET OPPORTUNITY

A young, fast-growing category with no established default.

The AI governance market was sized at roughly \$0.89B in 2024, growing at about 45.3% CAGR (MarketsandMarkets). The category is being defined now — incumbents are late to AI, and AI-native players are early. The opportunity is not to win a share of an existing market, but to define and own an emerging one.

08 BUSINESS MODEL & PRICING

Recurring SaaS, priced by AI risk, grown through expansion.

	Starter	Growth	Enterprise
Price	\$500 / mo	\$1,000 / mo	From ~\$3–5K / mo
Billed annually	\$6,000 / yr	\$12,000 / yr	\$40K–120K+ / yr
Frameworks	1–2	All frameworks	All + custom
AI systems	up to ~5	up to ~25	Unlimited
Best for	Seed–Series A	Series A–B (sweet spot)	Regulated / multi-framework

Beyond the subscription

Revenue compounds through audit-readiness sprints, policy & artifact generation, the Verified badge, usage overage, and — as the dataset matures — benchmark data. Compliance is never one-time, so every new AI system and every new regulation is a reason to pay more, at near-zero new cost to serve.

09 UNIT ECONOMICS

Defensible even conservatively; the upside is engineered, not assumed.

Metric	Conservative	Target
Blended ACV	\$10,000	\$20,000
Gross margin	70%	82%
LTV : CAC	~7.8 : 1	~13–15 : 1
CAC payback	~9 months	~4.5 months
Net revenue retention	—	110–120%

Illustrative, calibrated to early data. Both states clear healthy-SaaS benchmarks (3:1, <12-month payback). The largest single lever is automating the regulatory-update engine, moving gross margin from ~70% toward 85%.

10 TWO-YEAR PLAN

Define and own the AI Trust category.

The 24-month objective is category ownership, not just growth: make the Trust Score a recognised signal, prove global reach, harden the data moat, and build an ecosystem. Four six-month phases — Foundation & Wedge, Proof & Repeatability, Scale & Standard, Category & Moat — each gated on a single exit milestone.

11 FIVE-YEAR VISION

Become foundational infrastructure the AI economy runs on.

Over five years, CompliVibe moves from category leader to standard to infrastructure: “CompliVibe Verified” as the mark buyers, insurers, and regulators look for; continuous, autonomous trust; the definitive AI-risk dataset; a seat at the table on regulation; and the trust rails for the emerging agent economy.

12 OPERATIONS

Three layers; for a trust company, operational rigour is the product.

- **Platform Ops** — secure, reliable, always-current platform (incl. the 48h update engine and our own SOC 2/ISO posture).
- **Customer Ops** — onboarding, success, support, renewals — driving retention and expansion.
- **Business Ops** — GTM, finance, legal, people, and a single metrics dashboard — run lean and compliant.

13 POLICIES & GOVERNANCE

We hold ourselves to the standard we sell.

A full internal policy set spans governance & ethics, information security, data protection & privacy, AI governance on our own systems, resilience & third-party risk, people, and a customer-facing legal layer (ToS, Privacy Policy, DPA, SLA, responsible disclosure). Being audited on our own standard is a proof point, not overhead.

14 MOAT & COMPETITION

We win a newer game the incumbents are late to.

- **Category timing** — we are AI-native; incumbents (Vanta, Drata) retrofit AI onto security compliance.
- **Data moat** — every assessment compounds a benchmark dataset competitors cannot copy.
- **IP** — 5 patents filed, including the Shadow AI Discovery Engine.
- **Speed** — 48-hour regulatory update velocity turns every rule change into an upsell.
- **Unified layer** — governance-only players (Credo AI, Arthur, Holistic AI) don't automate compliance; we unify both.

15 TEAM

A product-first founding team with enterprise depth and a dedicated growth engine.

- **Adarsh Sharma** — Founder & CEO; the founding insight originated at Databricks.
- **Saransh Jha** — Co-founder, COO & CFO; 12 years of enterprise operations experience.
- **Rudra** — CMO, leading marketing and growth via the Cosmo Advisory partnership.
- **Om Panigrahi** — Growth Manager, reporting to the CEO, leading the dedicated growth team.

Growth is powered by an exclusive partnership with Cosmo Advisory, which dedicates a 10–12 person team to CompliVibe under a performance-based structure.

16 THE ASK & NEXT STEPS

Fuel to reach the proof point — repeatable pipeline and rising margin.

CompliVibe is raising a pre-seed round to accelerate go-to-market and automate the regulatory-update engine — the highest-leverage action in the model. Capital takes the company to the Phase 2 proof point: a repeatable global pipeline and gross margin trending toward target, ahead of the EU AI Act forcing function.

Round terms and the current data room are available on request at dataroom.complivibe.in.

We are not building a compliance tool. We are building the trust layer for the AI era — global from day one, and inevitable as AI becomes foundational to every company on Earth.